

Lab 8 Regression

Chen Wang, Junke Wang, Zhuozhi Xiong

due 11/27/19

Together as a class:

We are first going to install and load packages that we will need.

1. Install and load both the `modelr` package and the `openintro` package. Load the `tidyverse` package.

For this problem we will be using the data set called `gpa` from the `openintro` package. We are interest in association between the number of hours of sleep a student gets and their `gpa`.

```
install.packages("modelr")
```

```
## Installing package into '/home/rstudio-user/R/x86_64-pc-linux-gnu-library/3.6'  
## (as 'lib' is unspecified)
```

```
install.packages("openintro")
```

```
## Installing package into '/home/rstudio-user/R/x86_64-pc-linux-gnu-library/3.6'  
## (as 'lib' is unspecified)
```

```
library(modelr)  
library(openintro)
```

```
## Please visit openintro.org for free statistics materials
```

```
##
```

```
## Attaching package: 'openintro'
```

```
## The following objects are masked from 'package:datasets':
```

```
##
```

```
##      cars, trees
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyv
```

```
## v ggplot2 3.2.1    v purrr   0.3.3
```

```
## v tibble  2.1.3    v dplyr  0.8.3
```

```
## v tidyr   1.0.0    v stringr 1.4.0
```

```
## v readr   1.3.1    v forcats 0.4.0
```


```
## -- Conflicts ----- tidyv
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()    masks stats::lag()
```

2. Make a scatter plot and describe the association that you see.

```
gpa %>%  
  ggplot(aes(x=sleepnight, y=gpa)) +  
  geom_point()
```



It doesn't seem that there is a relationship between hours of sleep and GPA. We can't tell if there is a positive or a negative trend.

3. Create the linear regression model for this relationship between hours of sleep and GPA. We are regressing gpa onto hours of sleep. Call this model gpa_model.

```
gpa_model <- lm(gpa ~ sleepnight, data=gpa)
```

```
gpa_model %>%  
  summary()
```

```
##  
## Call:  
## lm(formula = gpa ~ sleepnight, data = gpa)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -0.67898 -0.22123  0.02102  0.21627  1.08110   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  3.46000    0.31819  10.874 4.14e-15 ***  
## sleepnight   0.01983    0.04458   0.445  0.658      
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 0.3381 on 53 degrees of freedom  
## Multiple R-squared:  0.003719,    Adjusted R-squared:  -0.01508   
## F-statistic: 0.1978 on 1 and 53 DF,  p-value: 0.6583
```

4. Find the correlation of this regression line and give the estimated β values.

```
correlation <- sqrt(0.003719) #add sign of intercept
correlation
```

```
## [1] 0.0609836
```

```
cor(gpa$gpa, gpa$sleepnight)
```

```
## [1] 0.06098308
```

the correlation is 0.061. The estimated intercept is 3.46. The estimated slope is 0.01983.

5. Interpret the $\hat{\beta}$ values. There are 2.

GPA = intercept + slope(hours of sleep)

We can interpret the intercept as us having an estimated GPA of 3.46 when we have 0 hours of sleep.

We can interpret the slope as for every 1 extra hour of sleep we get, we will increase our GPA by 0.01983 points.


6. Add the predicted values of GPA and the Residuals to the data frame `gpa` using the `add_predictions()` and `add_residuals()` functions from the `modelr` package.

```
gpa <- gpa %>%  
  add_predictions(gpa_model) %>%  
  add_residuals(gpa_model)
```


7. Now we want to check the conditions needed to use the least squares regression line. Create a qqplot and a residual plots in order to check the conditions. Are the conditions meet? Are there any outliers?
- 1 - independence of data
 - 2 - linear relationship of GPA and hours of sleep

4 - constant variability

```
gpa %>%  
  ggplot(aes(sample=resid))+  
  geom_qq()
```



```
gpa %>%  
  ggplot(aes(x=sleepnight,y=resid))+  
  geom_point()+  
  geom_hline(yintercept=0)
```



All of our assumptions check out. We can use linear regression.

8. Conduct a hypothesis test to see if there is an association between how much sleep a student gets and their GPA. What can we conclude?

$$H_0 : \beta_1 = 0$$

(We assume the slope is 0. Hours of sleep and GPA have no linear relationship.)

$$H_0 : \beta_1 \neq 0$$

(We want to prove that the slope is not 0. That hours of sleep and GPA have a linear relationship of some kind.)

```
gpa_model %>%  
  summary()
```

```
##  
## Call:  
## lm(formula = gpa ~ sleepnight, data = gpa)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -0.67898 -0.22123  0.02102  0.21627  1.08110   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  3.46000    0.31819  10.874 4.14e-15 ***  
## sleepnight   0.01983    0.04458   0.445  0.658      
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 0.3381 on 53 degrees of freedom  
## Multiple R-squared:  0.003719,    Adjusted R-squared:  -0.01508   
## F-statistic: 0.1978 on 1 and 53 DF,  p-value: 0.6583
```

Our t-statistics is 0.445, our p-value is 0.658. Because our p-value is greater than 0.05, we fail to reject the null hypothesis, we cannot say there is a linear relationship between hours of sleep and GPA.

On your own:

In this problem, we will be using the `babies` data set from the `openintro` package. Information about the data set can be found by running the following code in an `R` chunk: `?babies`. We are interested in finding out how the length of gestation (`gestation`), the age of the mother (`age`), the height of the mother (`height`), the mother's weight (`weight`), and whether or not this was the mother's first pregnancy (`parity`) are related to how much the baby will weight at birth (`bwt`).

1. Load the data and save it as `baby`.

```
baby <- babies
```

2. Create the linear regression model that regresses `bwt` the variables mentioned above. Print out the summary.

```
baby_model <- lm(bwt ~ gestation + age + height + weight + parity, data= baby)
```

```
baby_model %>%  
  summary()
```

```
##  
## Call:  
## lm(formula = bwt ~ gestation + age + height + weight + parity,  
##     data = baby)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -54.569 -10.506   0.453  10.063  54.285   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept) -86.59435    14.75733  -5.868 5.73e-09 ***  
## gestation    0.46579     0.02992  15.568 < 2e-16 ***  
## age          0.05233     0.08820   0.593  0.5531      
## height       1.04614     0.21064   4.967 7.82e-07 ***  
## weight       0.06564     0.02584   2.540  0.0112 *     
## parity      -2.96394     1.16403  -2.546  0.0110 *     
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 16.35 on 1178 degrees of freedom  
##    (52 observations deleted due to missingness)  
## Multiple R-squared:  0.2106, Adjusted R-squared:  0.2072   
## F-statistic: 62.84 on 5 and 1178 DF,  p-value: < 2.2e-16
```

3. Interpret the coefficients for `gestation` and `parity`.

The coefficient for `gestation` is 0.46579, which means when other explanatory variables don't change, and only `gestation` increases by 1 unit, the weight of baby at birth will increase by 0.46579 unit.

The coefficient for `parity` is -2.96394, which means when other explanatory variables don't change, and only `parity` increases by 1 unit, the weight of baby at birth will decrease by 2.96394 units.

4. Does the intercept value have any meaning in context?

Theoretically, the intercept value means when `gestation`, `age`, `height`, `weight` and `parity` are all 0, then the weight of the baby will -86,59435. But since these factors can not be 0, and baby's weight can not be negative, so the intercept value doesn't have meaning in context.

5. What is the Multiple R^2 value for this model.

```
baby_model %>%  
  summary()
```

```
##  
## Call:  
## lm(formula = bwt ~ gestation + age + height + weight + parity,  
##     data = baby)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -54.569 -10.506   0.453  10.063  54.285   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept) -86.59435    14.75733  -5.868 5.73e-09 ***  
## gestation     0.46579     0.02992  15.568 < 2e-16 ***  
## age           0.05233     0.08820   0.593  0.5531      
## height        1.04614     0.21064   4.967 7.82e-07 ***  
## weight        0.06564     0.02584   2.540  0.0112 *     
## parity       -2.96394     1.16403  -2.546  0.0110 *     
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 16.35 on 1178 degrees of freedom  
## (52 observations deleted due to missingness)  
## Multiple R-squared:  0.2106, Adjusted R-squared:  0.2072   
## F-statistic: 62.84 on 5 and 1178 DF,  p-value: < 2.2e-16
```

We can see that multiple R^2 for this model is 0.2106.

6. Add the residuals and predicted values to the data frame.

```
baby <- baby %>%  
  add_predictions(baby_model) %>%  
  add_residuals(baby_model)
```


7. Create a residual plot and a qqplot. Comment on whether or not the conditions are met to use the model you found in part 2.

The conditions include:

- 1 - independence of data
- 2 - linear relationship
- 3 - normality of residuals
- 4 - constant variability


```
baby %>%  
  ggplot(aes(sample=resid))+  
  geom_qq()
```

```
## Warning: Removed 52 rows containing non-finite values (stat_qq).
```




```
baby %>%  
  ggplot(aes(x=gestation,y=resid))+  
  geom_point()+  
  geom_hline(yintercept=0)
```

```
## Warning: Removed 52 rows containing missing values (geom_point).
```




```
baby %>%  
  ggplot(aes(x=age,y=resid))+  
  geom_point()+  
  geom_hline(yintercept=0)
```

```
## Warning: Removed 52 rows containing missing values (geom_point).
```




```
baby %>%  
  ggplot(aes(x=height,y=resid))+  
  geom_point()+  
  geom_hline(yintercept=0)
```

```
## Warning: Removed 52 rows containing missing values (geom_point).
```




```
baby %>%  
  ggplot(aes(x=weight,y=resid))+  
  geom_point()+  
  geom_hline(yintercept=0)
```

```
## Warning: Removed 52 rows containing missing values (geom_point).
```



```
baby %>%  
  ggplot(aes(x=factor(parity),y=resid))+  
  geom_boxplot()+  
  xlab("parity")
```

```
## Warning: Removed 52 rows containing non-finite values (stat_boxplot).
```



Therefore, from the diagrams above, we can see that all the conditions are met: from the qqplot we can see the residual nearly follows normality and from residual plot we can see the variability is constant, so we can use a linear regression.

8. Obtain a 95% confidence interval for the coefficient on `gestation` and `age`. Interpret both confidence intervals.

```
baby_model %>%
  summary()

##
## Call:
## lm(formula = bwt ~ gestation + age + height + weight + parity,
##     data = baby)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -54.569 -10.506   0.453  10.063  54.285
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -86.59435   14.75733  -5.868 5.73e-09 ***
## gestation     0.46579    0.02992  15.568 < 2e-16 ***
## age           0.05233    0.08820   0.593  0.5531
## height        1.04614    0.21064   4.967 7.82e-07 ***
## weight        0.06564    0.02584   2.540  0.0112 *
## parity       -2.96394    1.16403  -2.546  0.0110 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 16.35 on 1178 degrees of freedom
## (52 observations deleted due to missingness)
## Multiple R-squared:  0.2106, Adjusted R-squared:  0.2072
## F-statistic: 62.84 on 5 and 1178 DF,  p-value: < 2.2e-16

z_cri <- qnorm(0.025)
l1 <- 0.46579 + z_cri*0.02992
r1 <- 0.46579 - z_cri*0.02992
l1

## [1] 0.4071479
r1

## [1] 0.5244321
l2 <- 0.05233 + z_cri*0.08820
r2 <- 0.05233 - z_cri*0.08820
l2

## [1] -0.1205388
r2

## [1] 0.2251988
```

The 95% confidence interval for coefficient on `gestation` is (0.4071479, 0.5244321), which means we are 95% confident that the coefficient on `gestation` is between 0.4071479 and 0.5244321.

The 95% confidence interval for coefficient on `age` is (-0.1205388, 0.2251988), which means we are 95% confident that the coefficient on `age` is between -0.1205388 and 0.2251988.